TRADE AND INVESTMENT · STRATEGIC ANALYSIS

SIMON LACEY | AUGUST 2026

The Terms of Entry

Europe responds to China's success in EVs by setting new terms for market access

Market Access X

The Custody Problem

Why Diversifying AI Vendors Answers the Wrong Question, and What the US–China Enclosure of the AI Stack Means for Firms Building Outside Both

On June 12, 2026, the US Commerce Department ordered Anthropic to stop supplying its two most capable AI models to any non-US person, and within hours the company disabled them for every customer on Earth. The action was novel in one respect that matters more than the headline: it was the first time the United States reached past chips and past companies to pull a deployed frontier model off the global market by directive. For a firm that had built a product on that model, the capability did not degrade or get more expensive. It vanished overnight. This report is about the risk that event exposed, why diversifying across AI vendors does almost nothing to address it, and what actually does.

This report is not investment, legal, or compliance advice. It is intended for professionals operating at the intersection of trade policy, geopolitics, and supply chain strategy.

Executive Summary

Both the United States and China now treat every layer of the AI stack as a strategic control surface, and they do so from opposite directions. The US restricts to deny capability to its rival; China restricts to keep its strategic assets and its people from leaving, while distributing models open-weight in a way that erodes the value of US restrictions. Firms building on AI sit underneath this contest, exposed to two distinct risks they generally tend to treat as one. This report separates these risks and shows why the conventional hedge fails against the one that just materialized.

- **The risk that materialized is yank-risk, not competition risk.** A state can withdraw access to an AI capability with little notice. On June 12 the US did exactly that to Anthropic's frontier models, for all foreign nationals, worldwide. Diversifying across vendors does not address this, because every US frontier lab sits behind the same jurisdiction and the same discretionary authority.
- **Enforceability does not simply fall as the chokepoint climbs the stack; it splits.** Closed, API-served models remain controllable, through the provider the state can compel directly, as the Anthropic directive proved. Open weights, once published, cannot be recalled by any state. The open–closed line within the model layer is where control is won or lost.
- **The yank-risk is administration-independent.** The restrictive prior administration tried to wall model weights in by rule and the rule was rescinded; the current administration sells advanced chips to Chinese firms and sues its own states to deregulate AI, and it still pulled a frontier model from allied users overnight. No political configuration of US policy removes the exposure, because the instrument is executive discretion itself.
- **The courts are a real check, but a slow and uneven one the buyer cannot invoke.** An American company can litigate a withdrawal, and may win. But correction takes months, the national-security framing is where courts defer most, and the foreign firm depending on the model is a bystander to a case it has no standing in. The order need not be lawful to cause the harm; it need only be issued.

- **The answer is capability custody, along a continuum, not vendor diversification.** For the workloads where a sudden cutoff would be existential, the hedge is to depend on a capability that no single state can withdraw from the market, which today means open weights. For most firms that means logical custody (open weights served by a third party, re-sourceable, zero capital), not buying GPUs. Physical self-hosting is a costly far end reserved for genuine sovereignty or residency drivers.

90 min	~4 months	$0	1st
Time Anthropic had to disable its frontier models after the June 12 directive	Public-benchmark lead of closed frontier models over the best open weights, and growing	Capital required for logical custody of open weights via a third-party host	Time the US pulled a deployed frontier model from the global market by directive

1. The Firm That Cannot Leave

In roughly one hundred days, two governments demonstrated the same proposition from opposite ends of the world. In late February and March 2026, the US Department of War designated Anthropic, an American company, a supply-chain risk to national security, and a federal judge in California blocked the designation as unconstitutional retaliation. In April, China ordered Meta to unwind its roughly two-billion-dollar acquisition of the Singapore-based AI startup Manus, and reporting indicates Beijing barred two of the company's founders from leaving the country while the review proceeded. In June, the Trump administration returned to Anthropic and pulled its two most capable models off the global market by export directive. Three interventions, two states, one shared premise: a frontier AI asset is not free to organize itself however it wishes.

The two states reach for that premise through different instruments, and the difference is instructive. China encloses by holding the person and the transaction; it kept Manus's capability inside its borders by reaching the founders and the deal. The United States encloses by holding the substrate; it does not need to detain anyone, because it controls the compute, the corporate entity, and now the model weights that a US lab depends on. China grabs the asset at the perimeter. The US grabs it at the foundation. Both produce the same result: the frontier lab discovers it is the least escapable strategic asset of the modern era.

This matters because the intuitive response to jurisdictional risk is to imagine an exit. If a US lab is exposed to US control, why not relocate? The answer is that a frontier lab cannot meaningfully expatriate, because the things that make it valuable are each independently bound to a jurisdiction. Its training runs sit on compute in identifiable data centers. Its capability lives in weights that the same June directive showed can be export-controlled. Its work depends on a small, visible pool of researchers whose nationality and location are known. Its corporate form, its capital, and its governance are creatures of the law it was incorporated under. A founder can board a plane; the lab cannot follow, because moving all four substrates at once is not something any single decision can accomplish, and the state needs to block only one to make the move worthless.

The orbital escape hatch closes the same way

The most literal version of the escape fantasy is to leave the planet. Proposals for solar-powered orbital data centers, including a filing to launch up to a million satellites, are sometimes read as a route around terrestrial control. They are the opposite. Reaching orbit requires a launch license, one of the most tightly controlled activities that exists. The hardware lofted is built from the same export-controlled accelerators whether it sits in Virginia or in low Earth orbit, and the operator remains a company under its home jurisdiction wherever the satellites fly. Space does not dissolve the substrate hooks; it adds a licensing layer on top of them. The one domain where the US has the longest record of controlling technology that physically leaves the planet is, after all, satellites.

One clarification is owed on the Anthropic case, because it is easy to flatten. Anthropic was targeted twice and the outcomes diverged. It won the constitutional argument against the first, procurement-based designation, where a court found the government had punished protected speech. It lost a parallel motion in a second court, which deferred to the government on national-security grounds, and the model-export action of June is newer still and unresolved. The accurate description is not a company hobbled, nor a company vindicated, but a company targeted, partially checked, and still contesting. That mixed record is itself the point developed below: the check on this kind of power is real but uneven, and it arrives late.

2. Where Control Splits: The Open–Closed Line

The natural way to describe the contest is to say that control gets harder as it moves up the stack, from chips, to manufacturing know-how, to the models themselves. That description is half right and, stated plainly, wrong in a way that a sophisticated observer will catch. Hardware controls are genuinely enforceable: chips are physical, customs-legible, and the foreign direct product rule extends US reach over them across borders. Controls on people and know-how are softer, harder to police, but real. The error is to assume the trend continues smoothly into the model layer, so that model controls are simply the weakest of all. They are not uniformly weak. Enforceability does not keep declining; it splits.

The split runs along the line between closed and open weights. A closed, API-served model is controllable, not because there is a robust standing rule for it, but because there is a company the state can compel directly. The June 12 directive is the proof: a letter on a Friday afternoon took a deployed frontier model off the global market by later the same day. Open weights are the opposite. Once a model's weights are published, they are copied, mirrored, and re-hosted beyond any single state's recall. There is no company to compel and no act of withdrawal that reaches what is already in the wild. The model layer therefore contains both the most discretionarily controllable form of AI capability and the only form that is effectively beyond control, separated by whether the weights were ever released.

The counter-case deserves a direct hearing. Closed-model control at the access layer is not trivial to dismiss: providers run know-your-customer checks, gate access by API, and apply tiered restrictions to their most capable systems, and self-hosting a frontier-scale model is itself compute-bound enough to act as a chokepoint. More importantly, open weights are not yet a clean substitute for the frontier. Independent trackers put the best open-weight models roughly four months behind the closed frontier on public benchmarks, with a wider gap on private ones, and that gap has been growing rather than closing since early 2025. The split is real, but it is a split between controllable-and-current capability and uncontrollable-but-trailing capability. That trade-off, not a simple ranking, is what a buyer is actually choosing between.

3. The Yank Is Administration-Independent

A foreign firm weighing dependence on US AI might reasonably hope to wait out the politics, reasoning that the exposure is an artifact of a particular administration's restrictiveness and will ease when the policy turns. The record of the eighteen months before the June directive forecloses that hope, because it shows the exposure surviving a complete reversal of policy posture.

The prior administration's instinct was to wall AI capability in. In January 2025 it issued a framework that would have imposed the first export controls on AI model weights and a worldwide licensing requirement on advanced chips. The current administration rescinded that framework in May 2025, on the eve of its compliance date, calling it innovation-stifling and diplomatically damaging, and it has not been replaced by a standing model-weight regime. Two months later, in July 2025, the same administration signed an executive order whose entire thrust was to promote the export of the American AI stack, to flood allied markets with US hardware, models, and standards so that American technology became the global default. By the start of 2026 it had gone further, loosening chip controls on China itself: it shifted advanced-chip licensing for China from presumptive denial to case-by-case review and cleared major Chinese firms to buy advanced US accelerators, subject to a government cut of the proceeds.

That is the most aggressively pro-diffusion, export-maximizing AI posture the United States has taken. And it is the posture under which the June directive happened. The same government selling advanced chips to Chinese companies, and suing its own states to strip away AI regulation, pulled an American frontier model off the global market and away from allied users overnight. The conclusion is not that the policy is incoherent, and the report does not allege hypocrisy; the two moves operate on different layers and against different targets. The conclusion is structural. The restrictive administration tried to constrain capability by rule, and the rule died. The permissive administration constrains capability by discretionary directive, and the directive worked instantly. The instrument that produces yank-risk is not any particular policy; it is the executive discretion that both poles of US policy preserve. A buyer cannot price that risk off the prevailing stance, because the stance and the risk move independently.

4. Why the Courts Do Not Solve This for the Buyer

The obvious objection is that this discretion is not unbounded, and the objection has real force. A court blocked the first designation against Anthropic in unusually direct language, and Anthropic's own public account of the June directive supplies the makings of a strong legal challenge: it says the government offered only verbal evidence of a narrow, non-universal jailbreak, amounting to little more than asking the model to read a codebase and identify flaws, and that the same capability is widely available from other models. If the controlled capability is freely available elsewhere, a control on one model does not serve the stated security objective, which is the core of a disproportionality or arbitrary-and-capricious argument. The discretion, in short, can be challenged, and sometimes the challenge will win.

That check is real but it does not rescue the buyer, for three reasons that compound. First, it is uneven. The same litigation that produced the constitutional win also produced a second ruling in which a different court deferred to the government, and export controls invoking national security are precisely the category where courts defer most. The speech-retaliation theory that prevailed against a procurement designation may not even be available against a Commerce licensing action. A correction is plausible; it is not assured, and the one adverse ruling on record points toward deference. Second, it is slow. The directive was effective the night it arrived; judicial correction of the earlier action took about a month to begin and the merits were still being argued months later. For a production system that went dark in June, vindication in the fourth quarter is not a remedy. Third, and most decisively, the buyer is a bystander. The party with standing to litigate is the lab, not the foreign firm that depends on it, and the firm has no control over the case, no ability to accelerate it, and no assurance that the lab's interests run with its own. A lab might settle on terms that restore domestic access while leaving the foreign-national restriction, which is exactly what the directive targeted, in place.

The legal theory is credible; the outcome is not predictable

That Anthropic has a plausible disproportionality and administrative-law argument against the June directive is an assessment of the legal theory, and a defensible one. It is not a prediction that the courts will vacate the order, and on the evidence available a prediction either way would be unfounded. The point for a firm making a sourcing decision does not depend on who wins. A plausibly unlawful directive still removed a frontier model from the global market for an indeterminate interval before any court could reach it. The exposure a buyer must architect against is that interval of lost access, not the eventual verdict.

5. The Precedent: When the US Has Done This Before

None of this is unprecedented in kind, only in subject. The United States has a documented history of applying security-grade export controls to a commercial technology and watching foreign buyers respond by engineering out their dependence on it. The clearest case is commercial satellites. After 1999, when jurisdiction over commercial communications satellites was moved to the munitions-control regime, foreign manufacturers began advertising their products as free of US content, and buyers began designing American components out of their systems specifically to avoid the licensing burden. The behavior is documented down to individual customers telling US vendors they would no longer buy American.

Use the mechanism, not the morality tale

The satellite case is often told as a clean story of permanent self-inflicted decline, and that telling is contested. The market-share figures cited for the episode come from different bases, exports, manufacturing revenue, and prime-contractor share, that should not be conflated, and at least one careful analysis shows US share recovering to a high level years later while still under the same controls. The durable lesson is therefore not that export controls permanently destroy an industry. It is narrower and uncontested: when the US applies security controls to a commercial technology, foreign buyers treat independence from US jurisdiction as a purchasing criterion and build around the controlled input. That buyer behavior is the precedent that matters here, and it does not depend on how the long-run market settled.

The rhyme with the present is exact. Independence from US jurisdiction was the feature foreign satellite buyers paid for after 1999; independence from US jurisdiction is the feature an open-weight model offers an AI buyer after June 2026. The mechanism is the same and the conclusion the buyer draws is the same: when the controlling state has shown it will reach for the control, the rational response is to reduce the dependence that makes the control bite. The open-weight field is, among other things, the market supplying exactly that independence, and its rapid adoption outside the US is consistent with buyers already making the satellite-era move.

6. What to Do: Custody, Not Diversification

The instinctive response to all of this is to diversify across vendors, and it is close to useless against the risk in question. Vendor diversification answers a competition problem: it protects against one supplier raising prices, degrading service, or failing commercially. It does nothing about a jurisdiction problem, because keeping three US frontier labs in the mix instead of one buys no protection when a single directive can reach all of them through the jurisdiction they share. What changed in June is that the jurisdiction problem became the live one, and it requires a different response.

That response is capability custody, and the distinction that makes it work is between two risks a firm usually conflates. Yank-risk is the danger that a state cuts off access to a capability. Compliance-risk is the danger that a state makes a firm's use of a capability unlawful. They pull in opposite directions. A closed model from a single jurisdiction carries high yank-risk and low compliance-risk; an open Chinese model carries low yank-risk and, for a firm with Western-facing exposure, potentially rising compliance-risk. The task is not to minimize one risk but to place each workload correctly given how existential a sudden cutoff would be and how much compliance exposure the firm carries.

Custody is best understood as a continuum, and the most important correction this report can offer is that it does not mean buying GPUs. At one end sits a closed model reached through a single jurisdiction's API: no custody, full exposure to a yank. In the middle, and this is the right answer for most firms, sits an open-weight model served by a third-party host. Because the weights are published and cannot be withdrawn from the market, the capability is yank-immune and the firm can move it between hosts or bring it in-house if forced. This is logical custody, and it requires no capital. Only at the far end sits physical custody, self-hosting open weights on owned hardware, which is genuinely expensive: a frontier-scale deployment runs into the high six figures in hardware alone, and the running economics rarely beat simply paying for an API. Physical custody is justified by a sovereignty, data-residency, or latency driver, not by cost, and it is a niche, not the recommendation. The property that does the work across the whole continuum is not ownership of hardware. It is dependence on a capability that no state can pull from the market.

Two caveats that govern the whole recommendation

First, on capability: custody removes yank-risk only for a capability that actually meets the need. On the specific dimension that triggered the June directive, offensive cybersecurity, open models currently trail the closed frontier on the available evidence, and no clean public benchmark yet shows parity. Capability diffuses to open weights within months, but a measurable gap persists today, so custody is a paid-for, capability-verified hedge for a given workload, not a free swap of one model for another. Second, on compliance: the legal status of Chinese open-weight models in Western-facing and government-facing supply chains is genuinely unsettled. The open-weight hedge that removes yank-risk can raise compliance-risk for a firm with US-facing exposure. This is a question for counsel against the firm's specific footprint, not one this report can resolve.

Forward-Looking Implications

The following distinguishes what is already in motion from what depends on variables that remain open. It is a horizon map, not a forecast.

Actor	Near term (2026)	Medium term (2027–2028)	Longer term
Non-US firm on a US frontier model	Inventory which workloads sit on closed single-jurisdiction APIs and which would be existential to lose. Treat the June directive as a live yank-risk, not a one-off.	Move existential workloads to logical custody of capability-verified open weights. Track whether the Mythos/Fable restriction lifts or hardens.	Yank-risk on US closed models persists across administrations; custody posture, not vendor count, determines resilience.
US policy posture	Export-promotion of the American stack coexists with discretionary national-security yank. No standing model-weight regime; control is ad hoc.	Watch whether a replacement model-weight rule emerges, and whether discretionary directives recur. The instrument is discretion, not a rule a buyer can track.	The exposure is administration-independent. A change of administration changes the route, not the existence, of the yank-risk.
Chinese open-weight field	Open models trail the closed frontier but diffuse capability within months and court international adoption with permissive licensing.	Capability gap on the controlled axes (notably cyber) is the variable to watch; parity there would change the custody calculus materially.	If the gap closes, open-weight custody becomes a near-complete hedge; if it widens, custody carries a real capability cost.
The affected buyer base	Foreign buyers begin pricing US-jurisdiction dependence as a risk, echoing the post-1999 satellite response.	Adoption of open and non-US capability as a hedge accelerates or stalls on the capability gap and on compliance clarity.	The market supplying jurisdiction-independence (open weights) is the structural beneficiary of every recurrence of the yank.

Actionable Takeaways

Each of the following is a concrete action, not a general orientation. The logic applies wherever in an organization these questions are being asked.

01 Separate yank-risk from compliance-risk before doing anything else

Establish, for each AI dependency, whether the danger you face is loss of access (yank-risk) or unlawful use (compliance-risk). They demand opposite hedges, and most firms manage them as one. A closed US model is mostly a yank-risk; a Chinese open model, for a Western-facing firm, is mostly a compliance-risk. This single distinction governs everything that follows.

02 Inventory which workloads sit on withdrawable capability

Map each AI dependency to where it sits on the custody continuum: closed single-jurisdiction API (fully withdrawable), open weights on a third-party host (yank-immune, re-sourceable), or self-hosted open weights. The June 12 directive showed that withdrawal can take effect in an afternoon, so the relevant question is not whether a provider is reliable but whether the capability can be pulled at all.

03 Move only the existential workloads to custody, and verify capability first

Reserve the custody move for workloads where a sudden cutoff would halt revenue or breach an obligation. For those, depend on a capability no state can withdraw, which today means open weights. But verify that the open model actually meets the need on your specific task before relying on it; on some axes a measurable gap to the closed frontier persists. Custody is a capability-verified hedge, not a free substitution.

04 Default to logical custody; treat self-hosting as a niche

For almost all firms, custody means open weights served by a third party, which is yank-immune and requires no capital. Self-hosting on owned hardware is expensive and rarely justified on cost; reserve it for a genuine sovereignty, data-residency, or latency driver. Do not let the word custody be heard as a mandate to build a datacenter.

05 Do not rely on the courts as your continuity plan

An American provider may litigate a withdrawal and may win, but the correction is slow, uneven, and not yours to invoke; you have no standing in the provider's case and no assurance its settlement protects your access. Architect against the interval of lost access, not against the eventual verdict. The order does not have to be lawful to take your capability offline.

06 Get the compliance question answered against your own footprint

The open-weight hedge that removes yank-risk can create compliance exposure if you have US-government-facing or Western-regulated business. The law here is unsettled and footprint-specific. Put the conflict-of-laws analysis to counsel before adopting a Chinese open model in a regulated supply chain, not after.

Indicators to Watch

These are operational triggers, each a specific observable that resolves a question this report leaves open.

01 Whether the Mythos/Fable restriction lifts or hardens

The administration signaled the lockdown might lift within weeks. If it lifts quickly, the episode still stands as proof of the yank mechanism and the precedent it sets; if it hardens or recurs, the yank-risk is confirmed as a standing feature, not an anomaly.

02 Any litigation ruling on the June directive

A merits ruling, or its absence, will show whether export controls invoking national security are as deference-protected as the earlier split suggested. Watch whether any court reaches the disproportionality argument that the controlled capability is available elsewhere.

03 The open-versus-closed capability gap on the controlled axes

Independent benchmarks on offensive-cyber and agentic capability are the measure of whether open-weight custody is a near-complete hedge or one that carries a real capability cost. A clean head-to-head closing the gap would change the calculus materially.

04 Whether a standing US model-weight rule replaces ad hoc directives

A replacement for the rescinded framework would make the yank-risk a trackable regime rather than a discretionary letter. Its absence means the risk stays unpredictable, which for a buyer is the worse condition.

Vendor diversification answers a competition problem; capability custody answers a jurisdiction problem. What changed on June 12, 2026 is that the jurisdiction problem became the live one, and the only capability a state cannot pull from the market is an open weight already in the wild. The check on that power is the courts, but the courts are slow, uneven, and not the buyer's to invoke. A firm building on AI outside the United States should architect against the interval of lost access, not against the verdict that may or may not follow it.

Sources & References

Axios, Bloomberg, Fortune, and TIME reporting on the Commerce Department directive to Anthropic regarding Fable 5 and Mythos 5 (June 12–14, 2026), including the operative licensing language and Anthropic's public statement on the directive and its rebuttal. Anthropic v. United States Department of War: preliminary injunction, US District Court, Northern District of California (Judge Rita F. Lin, March 26, 2026); D.C. Circuit denial of stay (April 8, 2026); merits oral argument (May 19, 2026). NPR, CNN, CNBC, and Financial Times reporting on China's block of the Meta–Manus acquisition (NDRC, April 27, 2026) and the reported exit restrictions on the Manus founders (March 2026). BIS Framework for Artificial Intelligence Diffusion (published January 13, 2025; ECCN 4E091) and its rescission (announced May 12–13, 2025). Executive Order 14320, Promoting the Export of the American AI Technology Stack (July 2025), and America's AI Action Plan. BIS final rule revising advanced-chip license review for China from presumption of denial to case-by-case (January 13–14, 2026); reporting on H200 clearances to Chinese firms (through May 2026). Epoch AI and the International AI Safety Report 2026 on the open-versus-closed capability gap. CyberGym (Wang et al., 2025) and Lyptus Research on offensive-cyber capability diffusion. MacroPolo / Paulson Institute Global AI Talent Tracker on the national origin and work location of top-tier AI researchers. CSIS Aerospace Security and related analyses on commercial-satellite export controls after 1999 and subsequent market-share recovery. SpaceX FCC filing on orbital satellite constellations and reporting on space-based data centers. Specific figures, dates, and quotations are drawn from the primary sources named above; several forward-looking statements are expressly marked in the text as assessments rather than established facts.

Disclaimer

This report is produced by Market Access X for professional information purposes only. It does not constitute investment advice, legal advice, or a recommendation to buy or sell any security or asset. The analysis reflects information available at the time of publication and is subject to change. Several claims are expressly marked as assessment or as unresolved; these are the author's analytical judgments, not established facts, and are flagged as such in the text. The legal status of particular sourcing and compliance choices is unsettled and firm-specific, and nothing here substitutes for advice of counsel against a firm's own facts. Market Access X accepts no liability for decisions made on the basis of this report.